

1. Sample 1 of water froze at 0°C .
2. Sample 2 of water froze at 0°C .
3. Sample 3 of water froze at 0°C .
-
- N. Sample N of water froze at 0°C .
- N+1. If all past samples of water froze at 0°C , then
this sample of water will freeze at 0°C .
-
- C. This sample of water will freeze at 0°C . (1-N+1)

Is this argument valid?

We already have an explanation of why we should believe premises 1-N.

What about premise $N+1$?

In our first reading, Hume addresses the question of whether we should believe premises like this by drawing a distinction between two different kinds of claims:

Premise N+1 appears to be like the claim that the sun will rise tomorrow:
it is a matter of fact rather than a matter of the relations of ideas, and so
cannot be known "by the mere operation of thought."

$N+1$. If all past samples of water froze at 0°C , then this sample of water will freeze at 0°C .

The Uniformity of Nature

The future will be like the past.

It seems as though, if we should believe in the Uniformity of Nature, we should believe $N+1$. So the basic question is whether we should believe in the Uniformity of Nature.

Is the negation of this claim a contradiction?

It seems not. So it seems that, if we should believe it, we must believe it
on the basis of experience.

After all, yesterday the future was like the past. And the same for the day before that. And this suggests an argument for the Uniformity of Nature.

1. Yesterday, the future was like the past.
2. The day before yesterday, the future was like the past.
3. The day before the day before yesterday, the future was like the past.

.....

N. N days ago, the future was like the past.

C. Today, the future will be like the past. (1-N)

It is hard to see how we could make the argument valid without adding a premise which was just a restatement of the very claim — the Uniformity of Nature — which we were trying to prove.

All massive objects
attract one another.

Every 24 hours, the earth
rotates on its axis.

It is very plausible that these are all claims which we should believe. But we have at present no explanation of this fact. After all, the only arguments we can give for these claims will be arguments like our argument that the sun will come up tomorrow. And these arguments are invalid, and hence not good arguments. This seems to show that

Foundationalism, as we have formulated it, is false.

Halley's comet will next be visible
from earth in 2061.

It also raises an important question. We think that at least some scientific theories are theories we should believe, and that other generalizations are not. Consider, for example,

People born under the sign of Taurus
are more likely to be stubborn than
people born under other signs.

This is a prediction of the theory of astrology. One shouldn't believe this prediction because one shouldn't believe the theory. But why not?

It is tempting to say something like: there's no good argument for astrology; there's no reason to think that it is true.

That leads to our central question: what does it mean for a theory to be **well-supported** by the evidence?

The examples we have already discussed give us a plausible answer to this question. Scientific theories typically involve certain generalizations. In the simplest case, they will be claims of the form

All F 's are G .

These are not claims which our senses can tell us to be true. But our senses can tell us that claims like this are true:

This particular
thing is an F ,
and it is G .

Let's call claims which are related in this way to generalizations **instances** of the
generalization.

Then we might say that a generalization is well-supported by the evidence just in case the following two conditions are satisfied:

Our senses tell us
that many instances
of the generalization
are true.

Our senses don't tell us that any instances of the generalization are false.

Induction $\rightarrow$ Belief

If you know many true instances of a generalization P , and don't know of any false instances of P , you should believe P .

I want to now look at two problems for Induction $\rightarrow$ Belief.

The first is called the paradox of the ravens. Consider the following

generalization:

Let's look at second problem for Foundationalism.

Suppose that I now begin to investigate my immediate environment for non-black things, checking whether they are ravens. I check my shirt, the lectern, my laptop — and it turns out that none of them are ravens.

Foundationalism says that you should discard any belief which is neither (i) certain nor (ii) based on sense experience or (iii) a conclusion of a good argument.

No Foundations → No Belief

If you can't be certain that P and your senses don't tell you that P and you can't give a good argument for P, you should not believe P.

Suppose that someone believes that Foundationalism is true. It appears that they should reason as follows: Foundationalism is true; so, I should continue to believe Foundationalism only if it is (i) certain or (ii) based on sense experience or (iii) a conclusion of a good argument; but in fact Foundationalism satisfies none of (i)–(iii); so, I should discard my belief in Foundationalism.

Foundationalism thus appears to be unstable, in the sense that it recommends that it itself not be believed. It is thus a bit like this sentence:

No one should believe me.

If this sentence is true, no one should believe it. If it is untrue, of course, no one should believe it. So, no one should believe it.

Parallel reasoning implies that no one should believe Foundationalism. So,
no one should be convinced by arguments, like our evidentialist arguments,
which use it as a premise.

The same reasoning can be used to object to two other negative rules of belief we have considered.

No Good Argument $\rightarrow$ No Belief

If you can't give a good
argument for P, don't believe P.

Doubt $\rightarrow$ No Belief

If you cannot rule out a situation which would make P false, you should not believe P.

Indeed, it looks like a general problem for many principles which say “Don’t believe P unless your belief satisfies a certain special condition!” For we can always ask whether the principle itself satisfies that special condition. And we have just seen that in many cases, it plausibly won’t.

We already have an explanation of why we should believe premises 1-N.

What about premise $N+1$?

believe premises like this by drawing a distinction between two different

In our first reading, Hume addresses the question of whether we should

kindest of claims:

it is a matter of fact rather than a matter of relation or ideas, and so

Premise $N+1$ appears to be like the claim that the sun will rise tomorrow:

cannot be known 'by the mere operation of thought.'

should believe $N+1$. So the basic question is whether we should believe

It seems as though, if we should believe in the Uniformity of Nature, we

in the Uniformity of Nature.

on the basis of experience.

It seems not. So it seems that, if we should believe it, we must believe it

After all, yesterday the future was like the past. And the same for the day

before that. And this suggests an argument for the Uniformity of Nature.

It is hard to see how we could make the argument valid without adding a

premise which was just a restatement of the very claim — the Uniformity

of Nature — which we were trying to prove.



All massive objects
attract one another.



Every 24 hours, the earth
rotates on its axis.

these arguments are invalid, and hence not good arguments. This seems to show that

It is very plausible that these are all claims which we should believe. But we have at

claims will be arguments like our argument that the sun will come up tomorrow. And

Foundationalism, as we have formulated it, is false.

present no explanation of this fact. After all, the only arguments we can give for the



Halley's comet will next be visible
from earth in 2061.

theories should believe, and that other generalizations are not. Consider, for example,

It also raises an important question. We think that at least some scientific theories are



People born under the sign of Taurus
are more likely to be stubborn than
people born under other signs.

because one shouldn't believe the theory. But why not?

This is a prediction of the theory of astrology. One shouldn't believe this prediction

reason to think that it is true.

It is tempting to say something like: there's no good argument for astrology; there's no

supported by the evidence?

That leads to our central question: what does it mean for a theory to be well-

Scientific theories typically involve certain generalizations. In the simplest case, they

visibility of the form

The examples we have already discussed give us plausible answers to this question.

All F 's are G .

us that ain't like this are true:

These are not claims which our senses can tell

This particular
thing is an F ,
and it is G .

Let's all claim which are related in this way to generalizations of the

generalization.

case the following two conditions are satisfied:

Then we might say that a generalization is well-supported by the evidence just in

Our senses tell us
that many instances
of the generalization
are true.

Our senses don't tell us that any instances of the generalization are false.



Induction $\rightarrow$ Belief

If you know many true instances of a generalization P , and don't know of any false instances of P , you should believe P .

The first is called the paradox of the ravens. Consider the following

generalization:

then are ravens.

Suppose that I now begin to investigate my immediate environment for non-black things, checking

whether they are raven's. I check my shirt, then, my laptop — and it turns out that none of

certain nor (ii) based on sense experience or (iii) a conclusion of a good

Foundationalism says that you should discard any belief which is neither (i)

argument.



No Foundations → No Belief

If you can't be certain that P and
your senses don't tell you that P
and you can't give a good
argument for P, you should not
believe P.

they should reason as follows: Foundationalism is true; so, I should continue

Foundationalisn.

experience or (iii) a conclusion of a good argument; but in fact

Suppose that someone believes that Foundationalism is true. It appears that

Foundationalism satisfied none of (i)–(iii); so, should discard my belief in

to believe Foundationalism only if it is (i) certain or (ii) based on sense

Foundationalism thus appears to be unstable, in the sense that it



recommends that it itself not be believed. It is thus a bit like this sentence:



No one should believe me.

If this sentence is true, then it should believe it. If it is untrue, of course, no



one should believe it. So, one should believe it.

none should be convinced by arguments, like evidentialist arguments,

Parallel reasoning implies that no one should believe Foundationalism. So,

which use it as a premise.

believe we have considered.

The same reasoning can be used to object to other negative rules of



argument for P , don't believe P .

No Good Argument $\rightarrow$ No Belief

If you can't give a good



false, you should not believe P.

If you cannot rule out a

situation which would make P



Indeed, it looks like a general problem for many principles which say ‘Don’t

always ask whether the principle itself satisfies that special condition. And

believe. But unless you believe a certain special condition! For we can

we have just seen that in many cases, it plausibly won't.